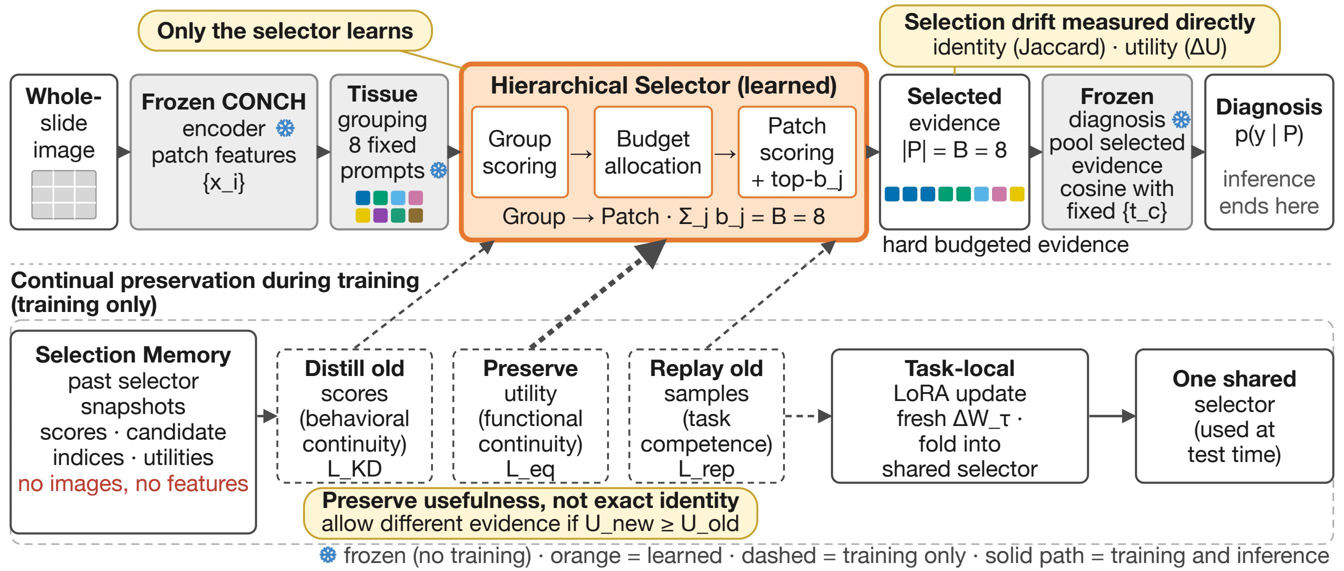
Only the selector learns
Selection drift measured directly
identity (Jaccard) · utility (ΔU)
Whole-slide image
Frozen CONCH encoder patch features {x_i}
Tissue grouping 8 fixed prompts
Hierarchical Selector (learned)
Group scoring
Budget allocation
Patch scoring + top-b_j
Group → Patch · Σ_j b_j = B = 8
Selected evidence |P| = B = 8
hard budgeted evidence
Frozen diagnosis pool selected evidence cosine with fixed {t_c}
Diagnosis p(y | P)
inference ends here
Continual preservation during training (training only)
Selection Memory
past selector snapshots
scores · candidate indices · utilities
no images, no features
Distill old scores (behavioral continuity) L_KD
Preserve utility (functional continuity) L_eq
Replay old samples (task competence) L_rep
Task-local LoRA update fresh ΔW_τ · fold into shared selector
One shared selector (used at test time)
Preserve usefulness, not exact identity
allow different evidence if U_new ≥ U_old
frozen (no training) · orange = learned · dashed = training only · solid path = training and inference